

The **Abstraction and Reasoning Corpus (ARC)** is a benchmark designed to evaluate general intelligence, not pattern memorization. Unlike typical vision tasks, ARC puzzles require humans to infer **abstract structural rules** — groupings, object relations, symmetry, color interactions — from a few input/output grid examples. Humans solve these tasks effortlessly, but AI systems struggle because ARC demands **conceptual reasoning**, not statistical learning.

The **ARGA** (Abstraction Reasoning Graph Analysis) framework approaches ARC by constructing **graph-based abstractions** and running a search-based solver over these abstractions. Each abstraction (e.g., ccg, nbccg, ccgbr, na) represents a different hypothesis about how to interpret objects and structures in the grid.

### Goal of the Project

We explore a simple but fundamental question:

**Can a Large Language Model (GPT-5) correctly predict which abstraction best describes an ARC task?**

If true, LLMs could guide ARC solvers by providing strong priors on *how to view* the task — an ability that humans rely on naturally.

We constructed two analytical datasets:

#### Dataset A — Solver Performance Across All Abstractions

For each task in our subset (54 tasks):

1. We modified ARGA so that it can be run with a manually selected abstraction.
2. We executed all nine abstractions independently on each task.
3. For each abstraction we recorded:
  - o Whether it solved the task
  - o Time taken (in seconds)
  - o Program length (length of apply\_call list)
4. We ranked the nine abstractions per task using:
  - o Time-based rank
  - o Program-length-based rank
    - using a 1-3-3-4 scheme (ties allowed below 1 second).

This gives the ground-truth notion of:

The “best” abstraction = fastest/shortest correct abstraction

## Dataset B — LLM-Predicted Abstraction Rankings

For the same tasks:

1. We crafted a detailed prompt describing:

- ARC
- ARGA
- All 9 abstractions
- JSON input/output format
- Task-specific reasoning guidelines

2. Prompt included:

- Persona priming
- Structural reasoning focus
- Explanations of object/group behaviors
- A self-reflection “mirroring” step

3. GPT-5 returned:

- Top-3 abstractions (predicted\_1, predicted\_2, predicted\_3)
- Justifications for each

We then converted these to rank vectors:

- predicted\_1 → rank 1
- predicted\_2 → rank 2
- predicted\_3 → rank 3
- all others → rank 9

Correlation Study:

Do GPT's Rankings Align With the Solver's Best Abstractions?

This is the central part of the project.

We compared, per task, the solver rank vector vs. the GPT-5 rank vector using:

Spearman's Rank Correlation Coefficient ( $\rho$ )

which measures monotonic agreement between rankings.

We compute two correlations:

1. LLM ranks vs solver time-based ranks
2. LLM ranks vs solver program-length ranks

### Correlation Results

Comparison	Spearman $\rho$	p-value	#Pairs
LLM vs Solver (Time)	0.275	1.03e-09	486
LLM vs Solver (Program Length)	0.272	1.11e-09	486

### Interpretation

A  $\rho \approx 0.27$  is:

- Weak in magnitude
- Highly statistically significant
- Meaningful for 9-class ranking tasks
- Far above random baselines ( $\sim 0$ )

### What the p-value Signifies

A p-value answers the question:

If the true correlation were zero (no relationship), how likely is it to observe a correlation as strong as ours just by random chance?

In our case:

- $p \approx 1 \times 10^{-9}$  (for both time and program length)
- This means the probability that our correlation happened by random noise is about: 0.000000001 (one in a billion).

Thus:

The correlation is extremely unlikely to be accidental.

GPT-5's abstraction rankings weakly align with the solver's performance-driven rankings.

Even though  $\rho \approx 0.27$  is moderate, the statistical certainty of the relationship is extremely high.

## Solver Top-3 Abstractions (Ground Truth)

Frequency of abstractions solving at least one task:

abstraction	solved_count
nbccg	38
ccgbr2	17
ccg	16
ccgbr	16
mcccg	15
lrg	15
nbehcg	11
nbvcg	10
na	7

From these:

## Solver Performance Summary

- nbccg alone solves 70.37%
- nbccg or ccgbr2 solves 77.78%
- nbccg, ccgbr2, or ccg solves 79.63%

## GPT-5 Abstraction Prediction Performance

Across all tasks:

Metric	Value
Top-1 accuracy	57.43%
Top-2 accuracy	83.33%
Top-3 accuracy	88.89%

Thus:

GPT-5 (88.89%) outperforms the solver's own 3 most frequent abstractions (79.63%).

## Qualitative Analysis: How GPT-5 Reasons

We studied rationales across predicted\_reason\_1/2/3.

Patterns include:

### **Strong Behaviors:**

- Detects connected components reliably
- Sensitive to background vs non-background
- Recognizes structural patterns like symmetry, alignment
- Applies abstraction rules correctly most of the time
- Differentiates pixel-level (na) from object-level abstractions

### **Weaknesses / Overgeneralization:**

- Sometimes over-prefers mcccg due to the presence of multi-color patches
- Slight bias toward high-level abstractions even when pixel-level ones work
- Occasionally describes transformations too literally

### **Conclusion:**

#### **This project shows that:**

1. GPT-5 can accurately infer the correct abstraction for an ARC task, with 88.89% success in its top-3 predictions.
2. These predictions correlate weakly with the abstractions that ARGA solves the fastest or with the shortest programs (Spearman  $p \approx 0.27$ ,  $p < 1e-9$ ).
3. GPT-5 outperforms ARGA's own empirically strongest top-3 abstractions.
4. The correlation suggests that:
  - GPT-5 may capture structural priors like in humans to some extent
  - GPT-5 could serve as an abstraction selector inside ARC solvers
  - This could somewhat reduce ARGA's search time